

# House Price Prediction Model using Machine Learning by Navjoth Singh

December 25, 2023

## 1 Data Collection: Importing Libraries & File

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: data = pd.read_csv("C:/Users/hp/OneDrive/Desktop/Data Science/Data for Practice/
↳HousingData.csv")
```

```
[3]: from sklearn.model_selection import train_test_split
```

```
[4]: from sklearn.linear_model import LinearRegression
```

```
[5]: from sklearn.metrics import mean_absolute_error, mean_squared_error
```

```
[6]: data.head()
```

```
[6]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.7	

	B	LSTAT	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	NaN	36.2

```
[7]: data.tail()
```

```
[7]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1	273	21.0	

502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273	21.0
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273	21.0
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273	21.0
505	0.04741	0.0	11.93	0.0	0.573	6.030	NaN	2.5050	1	273	21.0

	B	LSTAT	MEDV
501	391.99	NaN	22.4
502	396.90	9.08	20.6
503	396.90	5.64	23.9
504	393.45	6.48	22.0
505	396.90	7.88	11.9

```
[8]: data.sample(5)
```

```
[8]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
191	NaN	45.0	3.44	0.0	0.437	6.739	30.8	6.4798	5	398	
358	5.20177	0.0	18.10	1.0	0.770	6.127	83.4	2.7227	24	666	
129	0.88125	0.0	21.89	0.0	0.624	5.637	94.7	1.9799	4	437	
88	0.05660	0.0	3.41	0.0	0.489	7.007	86.3	3.4217	2	270	
106	0.17120	0.0	8.56	0.0	0.520	5.836	91.9	2.2110	5	384	

	PTRATIO	B	LSTAT	MEDV
191	15.2	389.71	4.69	30.5
358	20.2	395.43	11.48	22.7
129	21.2	396.90	18.34	14.3
88	17.8	396.90	5.50	23.6
106	20.9	395.67	18.66	19.5

```
[9]: data.shape
```

```
[9]: (506, 14)
```

```
[10]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CRIM        486 non-null    float64
1   ZN          486 non-null    float64
2   INDUS       486 non-null    float64
3   CHAS        486 non-null    float64
4   NOX         506 non-null    float64
5   RM          506 non-null    float64
6   AGE         486 non-null    float64
7   DIS         506 non-null    float64
8   RAD         506 non-null    int64
```

```

9    TAX      506 non-null    int64
10   PTRATIO  506 non-null    float64
11   B        506 non-null    float64
12   LSTAT    486 non-null    float64
13   MEDV     506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB

```

```
[11]: data.isnull().sum()
```

```

[11]: CRIM      20
      ZN        20
      INDUS    20
      CHAS      20
      NOX       0
      RM        0
      AGE      20
      DIS       0
      RAD       0
      TAX       0
      PTRATIO   0
      B         0
      LSTAT    20
      MEDV      0
      dtype: int64

```

```
[12]: data.nunique()
```

```

[12]: CRIM      484
      ZN        26
      INDUS    76
      CHAS      2
      NOX      81
      RM      446
      AGE     348
      DIS     412
      RAD      9
      TAX     66
      PTRATIO  46
      B      357
      LSTAT   438
      MEDV    229
      dtype: int64

```

```

[13]: data['CRIM'].fillna(data['CRIM'].mean(),inplace=True)
      data['AGE'].fillna(data['AGE'].mean(),inplace=True)

```

```
[14]: data['ZN'].fillna(data['ZN'].mean(),inplace=True)
data['INDUS'].fillna(data['INDUS'].mean(),inplace=True)
data['CHAS'].fillna(data['CHAS'].mean(),inplace=True)
data['LSTAT'].fillna(data['LSTAT'].mean(),inplace=True)
```

```
[15]: data.isnull().sum()
```

```
[15]: CRIM      0
      ZN       0
      INDUS   0
      CHAS    0
      NOX     0
      RM      0
      AGE     0
      DIS     0
      RAD     0
      TAX     0
      PTRATIO 0
      B       0
      LSTAT   0
      MEDV    0
      dtype: int64
```

```
[16]: data.describe()
```

```
[16]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	3.611874	11.211934	11.083992	0.069959	0.554695	6.284634	
std	8.545770	22.921051	6.699165	0.250233	0.115878	0.702617	
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	
25%	0.083235	0.000000	5.190000	0.000000	0.449000	5.885500	
50%	0.290250	0.000000	9.900000	0.000000	0.538000	6.208500	
75%	3.611874	11.211934	18.100000	0.000000	0.624000	6.623500	
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	

	AGE	DIS	RAD	TAX	PTRATIO	B	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	68.518519	3.795043	9.549407	408.237154	18.455534	356.674032	
std	27.439466	2.105710	8.707259	168.537116	2.164946	91.294864	
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000	
25%	45.925000	2.100175	4.000000	279.000000	17.400000	375.377500	
50%	74.450000	3.207450	5.000000	330.000000	19.050000	391.440000	
75%	93.575000	5.188425	24.000000	666.000000	20.200000	396.225000	
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000	


  

	LSTAT	MEDV
count	506.000000	506.000000

mean	12.715432	22.532806
std	7.012739	9.197104
min	1.730000	5.000000
25%	7.230000	17.025000
50%	11.995000	21.200000
75%	16.570000	25.000000
max	37.970000	50.000000

## 2 Exploratory Data Analysis (EDA)

```
[17]: data['AGE'].hist(bins=10, figsize = [10,5])
plt.grid(False)
plt.show()
```



```
[18]: data.columns
```

```
[18]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',
          'PTRATIO', 'B', 'LSTAT', 'MEDV'],
          dtype='object')
```

```
[19]: rows=2
cols=7

fig, axes = plt.subplots(rows, cols, figsize=(16, 5))


col = data.columns
index = 0
```

```

for i in range(rows):
    for j in range(cols):
        axes[i][j].hist(data[col[index]])
        axes[i][j].set_title(col[index])
        index = index + 1

plt.tight_layout()
plt.show()

```



```
[20]: data.corr()
```

```

[20]:
      CRIM      ZN      INDUS      CHAS      NOX      RM      AGE  \
CRIM      1.000000 -0.182930  0.391161 -0.052223  0.410377 -0.215434  0.344934
ZN      -0.182930  1.000000 -0.513336 -0.036147 -0.502287  0.316550 -0.541274
INDUS    0.391161 -0.513336  1.000000  0.058035  0.740965 -0.381457  0.614592
CHAS    -0.052223 -0.036147  0.058035  1.000000  0.073286  0.102284  0.075206
NOX      0.410377 -0.502287  0.740965  0.073286  1.000000 -0.302188  0.711461
RM      -0.215434  0.316550 -0.381457  0.102284 -0.302188  1.000000 -0.241351
AGE      0.344934 -0.541274  0.614592  0.075206  0.711461 -0.241351  1.000000
DIS     -0.366523  0.638388 -0.699639 -0.091680 -0.769230  0.205246 -0.724353
RAD      0.608886 -0.306316  0.593176  0.001425  0.611441 -0.209847  0.449989
TAX      0.566528 -0.308334  0.716062 -0.031483  0.668023 -0.292048  0.500589
PTRATIO  0.273384 -0.403085  0.384806 -0.109310  0.188933 -0.355501  0.262723
B       -0.370163  0.167431 -0.354597  0.050055 -0.380051  0.128069 -0.265282
LSTAT    0.434044 -0.407549  0.567354 -0.046166  0.572379 -0.602962  0.574893
MEDV    -0.379695  0.365943 -0.478657  0.179882 -0.427321  0.695360 -0.380223

      DIS      RAD      TAX      PTRATIO      B      LSTAT      MEDV
CRIM   -0.366523  0.608886  0.566528  0.273384 -0.370163  0.434044 -0.379695
ZN      0.638388 -0.306316 -0.308334 -0.403085  0.167431 -0.407549  0.365943
INDUS  -0.699639  0.593176  0.716062  0.384806 -0.354597  0.567354 -0.478657
CHAS   -0.091680  0.001425 -0.031483 -0.109310  0.050055 -0.046166  0.179882
NOX    -0.769230  0.611441  0.668023  0.188933 -0.380051  0.572379 -0.427321

```

RM	0.205246	-0.209847	-0.292048	-0.355501	0.128069	-0.602962	0.695360
AGE	-0.724353	0.449989	0.500589	0.262723	-0.265282	0.574893	-0.380223
DIS	1.000000	-0.494588	-0.534432	-0.232471	0.291512	-0.483429	0.249929
RAD	-0.494588	1.000000	0.910228	0.464741	-0.444413	0.468440	-0.381626
TAX	-0.534432	0.910228	1.000000	0.460853	-0.441808	0.524545	-0.468536
PTRATIO	-0.232471	0.464741	0.460853	1.000000	-0.177383	0.373343	-0.507787
B	0.291512	-0.444413	-0.441808	-0.177383	1.000000	-0.368886	0.333461
LSTAT	-0.483429	0.468440	0.524545	0.373343	-0.368886	1.000000	-0.721975
MEDV	0.249929	-0.381626	-0.468536	-0.507787	0.333461	-0.721975	1.000000

```
[21]: plt.subplots(figsize=(18,10))
sns.heatmap(data.corr(),annot=True,annot_kws={'size':14})
plt.show()
```



### 3 Train Test Split & Model Training

```
[22]: X = data.drop(['MEDV'], axis=1)
X.head()
```

```
[22]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	

```
4 0.06905 0.0 2.18 0.0 0.458 7.147 54.2 6.0622 3 222 18.7
```

```
      B      LSTAT
0 396.90  4.980000
1 396.90  9.140000
2 392.83  4.030000
3 394.63  2.940000
4 396.90 12.715432
```

```
[23]: y = data['MEDV']
      y.head()
```

```
[23]: 0    24.0
      1    21.6
      2    34.7
      3    33.4
      4    36.2
      Name: MEDV, dtype: float64
```

```
[24]: X_train, X_test, y_train, y_test = train_test_split (X, y, test_size=0.2,
      ↪random_state=0)

# The line of code you've mentioned performs a train-test split on your dataset,
↪using the train_test_split function from the scikit-learn library.
# Let's break down what each part of this line does:

# X: This typically represents your feature data, the input variables that your
↪model will learn from.
# y: This usually represents the target or output variable you want to predict,
↪based on your features.

# test_size=0.2: This parameter specifies that 20% of the data will be
↪allocated for testing, while 80% will be used for training.
# Adjusting this value changes the proportion of data used for training and
↪testing.

# random_state=0: This parameter sets the seed for the random number generator,
↪ensuring that the split is reproducible.
# Setting a specific random_state ensures that every time you run this code,
↪the data split will be the same.

# X_train: This variable stores the training data for the features (X) after
↪the split.
# This is the portion of your feature data that the model will learn from.

# X_test: This variable stores the testing data for the features (X) after the
↪split.
```



```
# This is the portion of your feature data that will be used to evaluate the
↪model's performance.

# y_train: This variable stores the training data for the target variable (y)
↪after the split.
# This corresponds to the target data associated with the X_train.

# y_test: This variable stores the testing data for the target variable (y)
↪after the split.
# This corresponds to the target data associated with the X_test.
```

```
[25]: X_train.shape, X_test.shape
      # 404 will undergo training and 506*20% = 102 will undergoes testing
```

```
[25]: ((404, 13), (102, 13))
```

```
[26]: y_train.shape, y_test.shape
```

```
[26]: ((404,), (102,))
```

```
[27]: model = LinearRegression()
      model.fit(X_train,y_train)
```

```
[27]: LinearRegression()
```

```
[28]: model.predict(X_test)
```

```
[28]: array([26.175296 , 22.64747588, 29.1456294 , 11.52971235, 21.65312134,
          19.42320699, 20.18413017, 21.46914355, 19.1985363 , 19.98228162,
           4.32483046, 16.16891668, 16.87682404,  5.31232373, 39.36827861,
          33.09358732, 21.9152876 , 36.61918436, 31.52676377, 23.52713482,
          24.96022461, 23.69866912, 20.88033802, 30.55074901, 22.74081741,
           8.66805959, 17.65119072, 17.93088633, 36.01223185, 21.16299556,
          17.83464361, 17.43306603, 19.5240167 , 23.50605522, 28.97262793,
          19.21808862, 11.23997435, 23.94256597, 17.86786717, 15.40849806,
          26.3630836 , 21.5193299 , 23.78733694, 14.84041522, 23.9445175 ,
          24.97067627, 20.11366175, 23.08636158, 10.42208266, 24.52832122,
          21.60847326, 18.66228165, 24.53362832, 31.03502944, 12.97457826,
          22.38536236, 21.34822822, 16.10928673, 12.37477824, 22.78596712,
          18.28714824, 21.91802045, 32.49771603, 31.21256855, 17.47867791,
          33.18861907, 19.17896285, 19.94662594, 20.17142015, 23.90228857,
          22.81288844, 24.17911208, 30.83402844, 28.87481037, 25.14581721,
           5.55072029, 37.0183454 , 24.15428003, 27.67587636, 19.63884644,
          28.74874123, 18.83204358, 17.63305678, 37.97947167, 39.49507972,
          24.17228966, 25.33605088, 16.75044819, 25.43224687, 16.65089426,
          16.49186628, 13.37283452, 24.81689254, 31.21188699, 22.0891919 ,
          20.49360168,  0.8229737 , 25.5004737 , 15.5481509 , 17.72901193,
```

```
25.77663998, 22.43131323])
```

```
[29]: y_predict = model.predict(X_test)
```

## 4 Regression Evaluation Metrics

```
[30]: from sklearn.metrics import r2_score
```

```
[31]: r2_score(y_test,y_predict)
```

```
[31]: 0.5703296053895559
```

```
[32]: mean_absolute_error(y_test,y_predict)
```

```
[32]: 3.9616211239591177
```


```
[33]: mean_squared_error(y_test,y_predict)
```

```
[33]: 34.987389544238766
```

```
[34]: np.sqrt(mean_squared_error(y_test,y_predict))  
# root_mean_squared_error
```

```
[34]: 5.915013909048631
```

```
[35]: plt.figure(figsize=(16,6))  
x_points = list(range(len(y_test)))  
plt.plot(x_points,y_test,label='y_true')  
plt.plot(x_points,y_predict,label='y_predict')  
plt.legend()  
plt.show()
```



## 5 Plot Learning Curve

```
[36]: from sklearn.model_selection import learning_curve, ShuffleSplit
```

```
[37]: def plot_learning_curves(estimator, title, X, y, ylim=None, cv=None,
    ↪train_size=np.linspace(0.1, 1.0, 10)):
    plt.figure()
    plt.title(title)
    plt.xlabel('Training Examples')
    plt.ylabel('Score')

    train_sizes, train_scores, test_scores = learning_curve(estimator, X, y,
    ↪train_sizes=train_size, cv=cv)

    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)

    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)


    plt.grid()

    plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                     train_scores_mean + train_scores_std, alpha=0.1,
    ↪color='red')
    plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.1,
    ↪color='green')

    plt.plot(train_sizes, train_scores_mean, 'o-', color='red', label='Training_
    ↪Scores')
    plt.plot(train_sizes, test_scores_mean, 'o-', color='green', label='Test_
    ↪Scores')

    plt.legend(loc='best')
    return plt
```


```
[38]: title = 'Learning Curve for Linear Regression'
cv = ShuffleSplit(n_splits=100, test_size=0.2, random_state=0)
model = LinearRegression()
plot_learning_curves(model, title, X, y, ylim=(0.7, 1.01), cv=cv)
plt.show()
```



## 6 Residuals Plot (Errors Plot)

```
[39]: from yellowbrick.regressor import ResidualsPlot
```

```
[40]: viz = ResidualsPlot(model)
viz.fit(X_train, y_train)
viz.score(X_test, y_test)
viz.show()
plt.show()
```




## 7 Prediction Error Plot

```
[41]: from yellowbrick.regressor import PredictionError
```

```
[42]: viz = PredictionError(model)
viz.fit(X_train, y_train)
viz.score(X_test, y_test)
viz.show()
plt.show()
```

C:\Users\hp\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names  
warnings.warn(



8 Thank You by Navjoth Singh